

IST438-W5-Applications

3/27/23

Decision trees

In this application, we will train decision trees for regression and classification tasks.

Packages

We need to install `{tidymodels}` package to train decision tree models. It is one of the most famous ML package in R because it consists many tools which are used in ML process:

- `{rsample}` is used to split dataset: `initial_split()`
- `{recipes}` for feature engineering
- `{parship}` model fitting
- `{tune}` model tuning
- `{yardstick}` model evaluation

Please use the two-step codes below: (1) install, (2) load the package.

```
#install.packages("tidymodels") # training models
#install.packages("DALEX").      # datasets
#install.packages("rpart.plot") # visualizing decision tree
library(tidymodels)
library(DALEX)
library(rpart.plot)
```

Dataset

`apartments` and `titanic` datasets are used in application to compare the performance of the regression models and decision trees.

Model training with {tidymodels} for regression task

Data splitting:

```
apartments_split <- initial_split(data = apartments, # dataset to split
                                   prop = 0.80)      # proportion of train set

apartments_train <- apartments_split |> training()
apartments_test  <- apartments_split |> testing()
```

Model specification:

- **type**: model type, e.g. regression, decision tree or etc.
- **engine**: different R packages have engines
- **mode**: learning task, e.g. regression or classification

Defining model specification:

```
dt_model <- decision_tree() |> # try linear_reg()
  set_engine("rpart") |>      # and lm
  set_mode("regression")
```

Model training:

```
dt_apartments <- dt_model |>
  fit(m2.price ~., data = apartments_train)
dt_apartments
```

parsnip model object

n= 800

node), split, n, deviance, yval
* denotes terminal node

- 1) root 800 658070200 3483.499
- 2) district=Bemowo,Bielany,Praga,Ursus,Ursynow,Wola 482 172944900 3009.589
- 4) surface>=87.5 229 49155870 2587.393
- 8) floor>=5.5 117 15143270 2313.538 *
- 9) floor< 5.5 112 16071740 2873.473 *
- 5) surface< 87.5 253 46022840 3391.735

```

10) floor>=4.5 148 17768190 3194.399 *
11) floor< 4.5 105 14367660 3669.886 *
3) district=Mokotow,Ochota,Srodmiescie,Zoliborz 318 212792000 4201.814
6) district=Mokotow,Ochota,Zoliborz 240 81098870 3884.808
12) surface>=74.5 144 28089320 3584.062
24) floor>=5.5 74 12182060 3366.878 *
25) floor< 5.5 70 8726800 3813.657 *
13) surface< 74.5 96 20448250 4335.927
26) floor>=5.5 59 6452186 4090.000 *
27) floor< 5.5 37 4737697 4728.081 *
7) district=Srodmiescie 78 33364750 5177.218
14) surface>=64.5 42 8690634 4722.476 *
15) surface< 64.5 36 5856217 5707.750 *

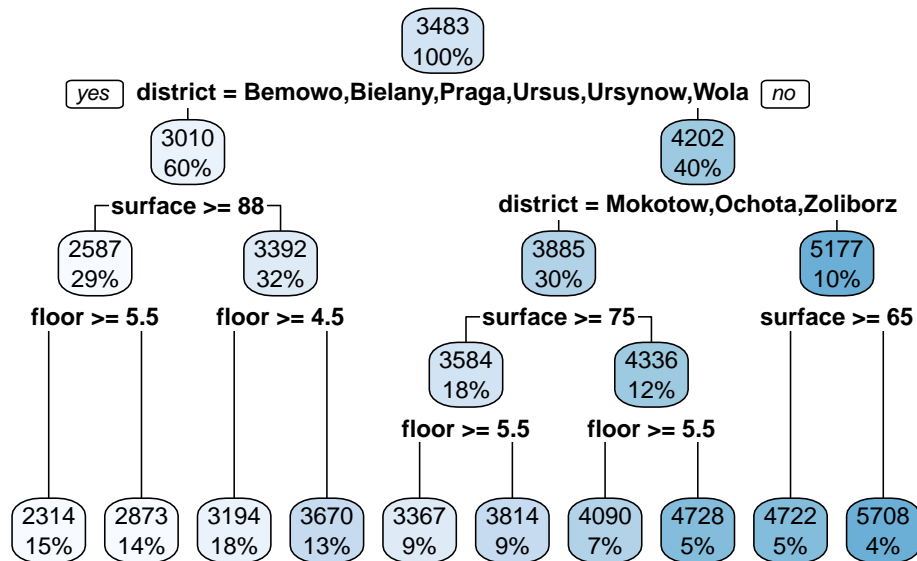
```

Visualizing the decision tree:

```
rpart.plot(dt_apartments$fit)
```

Warning: Cannot retrieve the data used to build the model (so cannot determine roundint and :
To silence this warning:

Call `rpart.plot` with `roundint=FALSE`,
or rebuild the `rpart` model with `model=TRUE`.



Make predictions by using the trained model:

```
apartments_predictions <- dt_apartments |>
  predict(new_data = apartments_test)

apartments_predictions
```

```
# A tibble: 200 x 1
  .pred
  <dbl>
1 3367.
2 2314.
3 4722.
4 4728.
5 3194.
6 4728.
7 3194.
8 3194.
9 2314.
10 4090
# ... with 190 more rows
```

Evaluating model performance:

{yardstick} package is used to evaluate/measure the model performance. Its functions require a data.frame or tibble with model results. To combine the model prediction and actual/observed values of target variable in test data, `cbind()` function can be used as follows:

```
apartments_results <- tibble(predicted = apartments_predictions$.pred,
                             actual    = apartments_test$m2.price)

apartments_results
```

```
# A tibble: 200 x 2
  predicted actual
  <dbl>    <dbl>
1    3367.    3517
2    2314.    2346
3    4722.    4745
4    4728.    3961
5    3194.    2797
6    4728.    5116
```

```

7      3194.    3172
8      3194.    3378
9      2314.    3372
10     4090     3868
# ... with 190 more rows

```

Then we can calculate the RMSE of the model by using `rmse()` function with obligatory arguments: `truth` must be assigned with actual values, `estimate` must be assigned with predicted values of target variable.

```
apartments_results |> rmse(truth = actual, estimate = predicted)
```

```

# A tibble: 1 x 3
  .metric .estimator .estimate
  <chr>   <chr>       <dbl>
1 rmse    standard      422.

```

R^2 metric can be calculated in similar manner by `rsq()` function:

```
apartments_results |> rsq(truth = actual, estimate = predicted)
```

```

# A tibble: 1 x 3
  .metric .estimator .estimate
  <chr>   <chr>       <dbl>
1 rsq     standard      0.784

```

Streaming model fitting by `last_fit()` function. It takes a model specification, model formula, and data split object.

```

apartments_last_fit <- dt_model |>
  last_fit(m2.price ~., split = apartments_split)

```

```
apartments_last_fit
```

```

# Resampling results
# Manual resampling
# A tibble: 1 x 6
  splits          id      .metrics .notes   .predictions .workflow
  <list>         <chr>    <list>  <list>  <list>       <list>
1 <split [800/200]> train/test split <tibble> <tibble> <tibble>    <workflow>

```

Collecting metrics:

```
apartments_last_fit |> collect_metrics()
```

A tibble: 2 x 4

	.metric	.estimator	.estimate	.config
	<chr>	<chr>	<dbl>	<chr>
1	rmse	standard	422.	Preprocessor1_Model1
2	rsq	standard	0.784	Preprocessor1_Model1

Collecting predictions:

```
apartments_last_fit |> collect_predictions()
```

A tibble: 200 x 5

	id		.pred	.row	m2.price	.config
	<chr>		<dbl>	<int>	<dbl>	<chr>
1	train/test	split	3367.	4	3517	Preprocessor1_Model1
2	train/test	split	2314.	8	2346	Preprocessor1_Model1
3	train/test	split	4722.	9	4745	Preprocessor1_Model1
4	train/test	split	4728.	11	3961	Preprocessor1_Model1
5	train/test	split	3194.	12	2797	Preprocessor1_Model1
6	train/test	split	4728.	13	5116	Preprocessor1_Model1
7	train/test	split	3194.	17	3172	Preprocessor1_Model1
8	train/test	split	3194.	18	3378	Preprocessor1_Model1
9	train/test	split	2314.	30	3372	Preprocessor1_Model1
10	train/test	split	4090	31	3868	Preprocessor1_Model1

... with 190 more rows

Model training with {tidymodels} for classification task

```
set.seed(123)
titanic_split <- initial_split(data = titanic, # dataset to split
                               prop = 0.80)   # proportion of train set

titanic_train <- titanic_split |> training()
titanic_test  <- titanic_split |> testing()
```

Model specification:

- **type**: model type, e.g. regression, decision tree or etc.
- **engine**: different R packages have engines
- **mode**: learning task, e.g. regression or classification

Defining model specification:

```
dt_model <- decision_tree() |>
  set_engine("rpart") |>
  set_mode("classification")
```

Model training:

```
dt_titanic <- dt_model |>
  fit(survived ~., data = titanic_train)

dt_titanic
```

parsnip model object

n= 1765

node), split, n, loss, yval, (yprob)
 * denotes terminal node

```
1) root 1765 590 no (0.66572238 0.33427762)
 2) gender=male 1362 302 no (0.77826725 0.22173275)
   4) class=2nd,3rd,engineering crew,restaurant staff,victualling crew 1165 212 no (0.8180
     8) age>=4.5 1141 196 no (0.82822086 0.17177914) *
     9) age< 4.5 24    8 yes (0.33333333 0.66666667) *
   5) class=1st,deck crew 197  90 no (0.54314721 0.45685279)
     10) age>=54.5 26    2 no (0.92307692 0.07692308) *
     11) age< 54.5 171  83 yes (0.48538012 0.51461988)
        22) fare>=39.5703 61  20 no (0.67213115 0.32786885) *
        23) fare< 39.5703 110  42 yes (0.38181818 0.61818182)
            46) country=Australia,Cuba,Mexico,Northern Ireland,Scotland,Wales 8    1 no (0.8750
            47) country=Belgium,Canada,Channel Islands,England,France,Germany,Ireland,Sweden,S
 3) gender=female 403 115 yes (0.28535980 0.71464020)
   6) class=3rd 184  87 no (0.52717391 0.47282609)
     12) fare>=23.07 30    3 no (0.90000000 0.10000000) *
     13) fare< 23.07 154  70 yes (0.45454545 0.54545455)
        26) country=Belgium,Croatia,Croatia (Modern),Sweden 27    6 no (0.77777778 0.22222222)
        27) country=Denmark,England,Finland,Ireland,Lebanon,Norway,Slovenia,Switzerland,Syria
```

```

54) age>=39.5 14 2 no (0.85714286 0.14285714) *
55) age< 39.5 113 37 yes (0.32743363 0.67256637) *
7) class=1st,2nd,restaurant staff,victualling crew 219 18 yes (0.08219178 0.91780822) *

```

Visualizing the decision tree:


```
rpart.plot(dt_titanic$fit)
```

Warning: Cannot retrieve the data used to build the model (so cannot determine roundint and
To silence this warning:

```

Call rpart.plot with roundint=FALSE,
or rebuild the rpart model with model=TRUE.

```



Make predictions by using the trained model. If you want to calculate the predicted probabilities, it is necessary to assign the `type` argument as "prob".

```

titanic_predictions <- dt_titanic |>
  predict(new_data = titanic_test)

titanic_predictions

```



```
# A tibble: 442 x 1
  .pred_class
  <fct>
1 no
2 yes
3 no
4 no
5 no
6 no
7 no
8 yes
9 no
10 no
# ... with 432 more rows
```

If you want to calculate the predicted probabilities, it is necessary to assign the `type` argument as "prob".

```
dt_titanic |>
  predict(new_data = titanic_test,
          type      = "prob")
```

```
# A tibble: 442 x 2
  .pred_no .pred_yes
  <dbl>    <dbl>
1  0.828    0.172
2  0.0822   0.918
3  0.828    0.172
4  0.828    0.172
5  0.828    0.172
6  0.828    0.172
7  0.828    0.172
8  0.0822   0.918
9  0.828    0.172
10 0.828    0.172
# ... with 432 more rows
```

Evaluating model performance:

`{yardstick}` package is used to evaluate/measure the model performance. Its functions require a data.frame or tibble with model results. To combine the model prediction and actual/observed values of target variable in test data, `cbind()` function can be used as follows:

```
titanic_results <- tibble(predicted = titanic_predictions$.pred_class,
                           actual    = titanic_test$survived)
```

Then we can calculate the RMSE of the model by using `rmse()` function with obligatory arguments: `truth` must be assigned with actual values, `estimate` must be assigned with predicted values of target variable.

```
titanic_results |> conf_mat(truth = actual,
                           estimate = predicted)
```

	Truth	
Prediction	no	yes
no	299	47
yes	22	74

The accuracy metric can be calculated in similar manner by `accuracy()` function:

```
titanic_results |> accuracy(truth = actual, estimate = predicted)
```

```
# A tibble: 1 x 3
  .metric .estimator .estimate
  <chr>   <chr>      <dbl>
1 accuracy binary      0.844
```

The sensitivity metric can be calculated in similar manner by `sens()` function:

```
titanic_results |> sens(truth = actual, estimate = predicted)
```

```
# A tibble: 1 x 3
  .metric .estimator .estimate
  <chr>   <chr>      <dbl>
1 sens   binary      0.931
```

```
titanic_results |> spec(truth = actual, estimate = predicted)
```

```
# A tibble: 1 x 3
  .metric .estimator .estimate
  <chr>   <chr>       <dbl>
1 spec    binary         0.612
```

{tidymodels} ecosystem provides many binary classification metrics:

- accuracy()
- kap()
- sens()
- spec()
- ppv()
- npv()
- mcc()
- j_index()
- bal_accuracy()
- detection_prevalence()
- precision()
- recall()
- f_meas()

Streaming model fitting by `last_fit()` function. It takes a model specification, model formula, and data split object.

```
titanic_last_fit <- dt_model |>
  last_fit(survived ~., split = titanic_split)
```

Collecting metrics:

```
titanic_last_fit |> collect_metrics()
```

```
# A tibble: 2 x 4
  .metric .estimator .estimate .config
  <chr>   <chr>       <dbl> <chr>
1 accuracy binary         0.844 Preprocessor1_Model1
2 roc_auc  binary         0.818 Preprocessor1_Model1
```

Collecting predictions:

```
titanic_last_fit |> collect_predictions()
```

```
# A tibble: 442 x 7
  id      .pred_no .pred_yes .row .pred_class survived .config
  <chr>      <dbl>    <dbl> <int> <fct>      <fct>    <chr>
1 train/test split  0.828    0.172     3 no        no      Preprocessor1~
2 train/test split  0.0822   0.918    21 yes      yes      Preprocessor1~
3 train/test split  0.828    0.172    22 no        no      Preprocessor1~
4 train/test split  0.828    0.172    28 no        no      Preprocessor1~
5 train/test split  0.828    0.172    42 no        no      Preprocessor1~
6 train/test split  0.828    0.172    43 no        no      Preprocessor1~
7 train/test split  0.828    0.172    47 no        no      Preprocessor1~
8 train/test split  0.0822   0.918    50 yes      yes      Preprocessor1~
9 train/test split  0.828    0.172    53 no        no      Preprocessor1~
10 train/test split 0.828    0.172    57 no        no      Preprocessor1~
# ... with 432 more rows
```

Model validation

```
set.seed(123)
titanic_folds <- vfold_cv(titanic_train,
                          v = 10)
titanic_folds
```

```
# 10-fold cross-validation
# A tibble: 10 x 2
  splits      id
  <list>     <chr>
1 <split [1588/177]> Fold01
2 <split [1588/177]> Fold02
3 <split [1588/177]> Fold03
4 <split [1588/177]> Fold04
5 <split [1588/177]> Fold05
6 <split [1589/176]> Fold06
7 <split [1589/176]> Fold07
8 <split [1589/176]> Fold08
9 <split [1589/176]> Fold09
10 <split [1589/176]> Fold10
```

```
titanic_wf <- workflow() |>
  add_model(dt_model) |>
```

```
add_formula(survived ~.)
```

```
titanic_fit_cv <- titanic_wf |>  
  fit_resamples(titanic_folds)
```

You can see the mean values of metrics over folds:

```
titanic_fit_cv |> collect_metrics()
```

```
# A tibble: 2 x 6
```

	.metric	.estimator	mean	n	std_err	.config
	<chr>	<chr>	<dbl>	<int>	<dbl>	<chr>
1	accuracy	binary	0.793	10	0.00672	Preprocessor1_Model11
2	roc_auc	binary	0.761	10	0.0110	Preprocessor1_Model11

Or you can check the metric values for each fold:

```
titanic_fit_cv |> collect_metrics(summarize = FALSE)
```

```
# A tibble: 20 x 5
```

	id	.metric	.estimator	.estimate	.config
	<chr>	<chr>	<chr>	<dbl>	<chr>
1	Fold01	accuracy	binary	0.774	Preprocessor1_Model11
2	Fold01	roc_auc	binary	0.688	Preprocessor1_Model11
3	Fold02	accuracy	binary	0.774	Preprocessor1_Model11
4	Fold02	roc_auc	binary	0.783	Preprocessor1_Model11
5	Fold03	accuracy	binary	0.808	Preprocessor1_Model11
6	Fold03	roc_auc	binary	0.775	Preprocessor1_Model11
7	Fold04	accuracy	binary	0.814	Preprocessor1_Model11
8	Fold04	roc_auc	binary	0.779	Preprocessor1_Model11
9	Fold05	accuracy	binary	0.797	Preprocessor1_Model11
10	Fold05	roc_auc	binary	0.798	Preprocessor1_Model11
11	Fold06	accuracy	binary	0.778	Preprocessor1_Model11
12	Fold06	roc_auc	binary	0.715	Preprocessor1_Model11
13	Fold07	accuracy	binary	0.818	Preprocessor1_Model11
14	Fold07	roc_auc	binary	0.776	Preprocessor1_Model11
15	Fold08	accuracy	binary	0.767	Preprocessor1_Model11
16	Fold08	roc_auc	binary	0.746	Preprocessor1_Model11
17	Fold09	accuracy	binary	0.778	Preprocessor1_Model11

18	Fold09	roc_auc	binary	0.783	Preprocessor1_Model11
19	Fold10	accuracy	binary	0.824	Preprocessor1_Model11
20	Fold10	roc_auc	binary	0.765	Preprocessor1_Model11